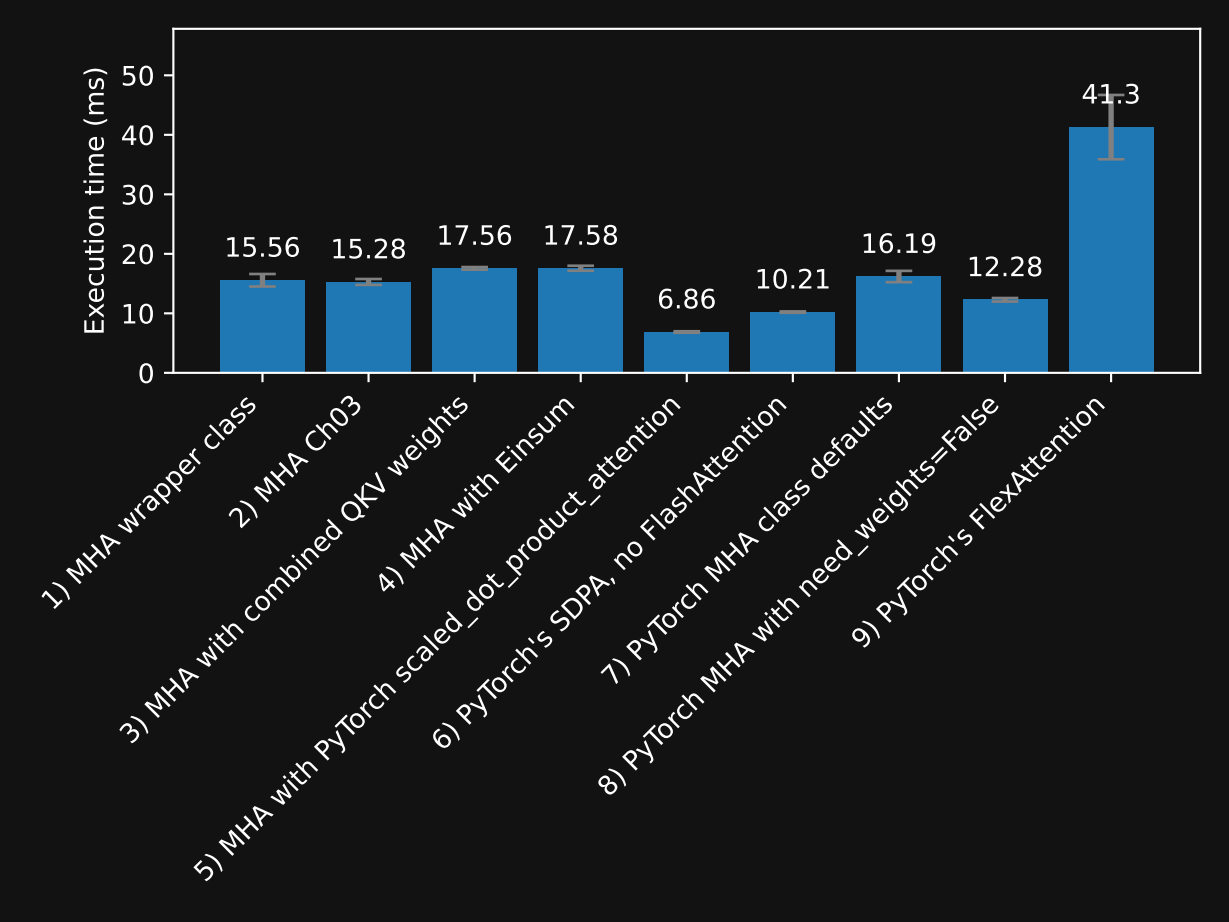
Execution time (ms)
0
10
20
30
40
50
15.56
15.28
17.56
17.58
6.86
10.21
16.19
12.28
41.3
1) MHA wrapper class
2) MHA Ch03
3) MHA with combined QKV weights
4) MHA with Einsum
5) MHA with PyTorch scaled_dot_product_attention
6) PyTorch's SDPA, no FlashAttention
7) PyTorch MHA class defaults
8) PyTorch MHA with need_weights=False
9) PyTorch's FlexAttention