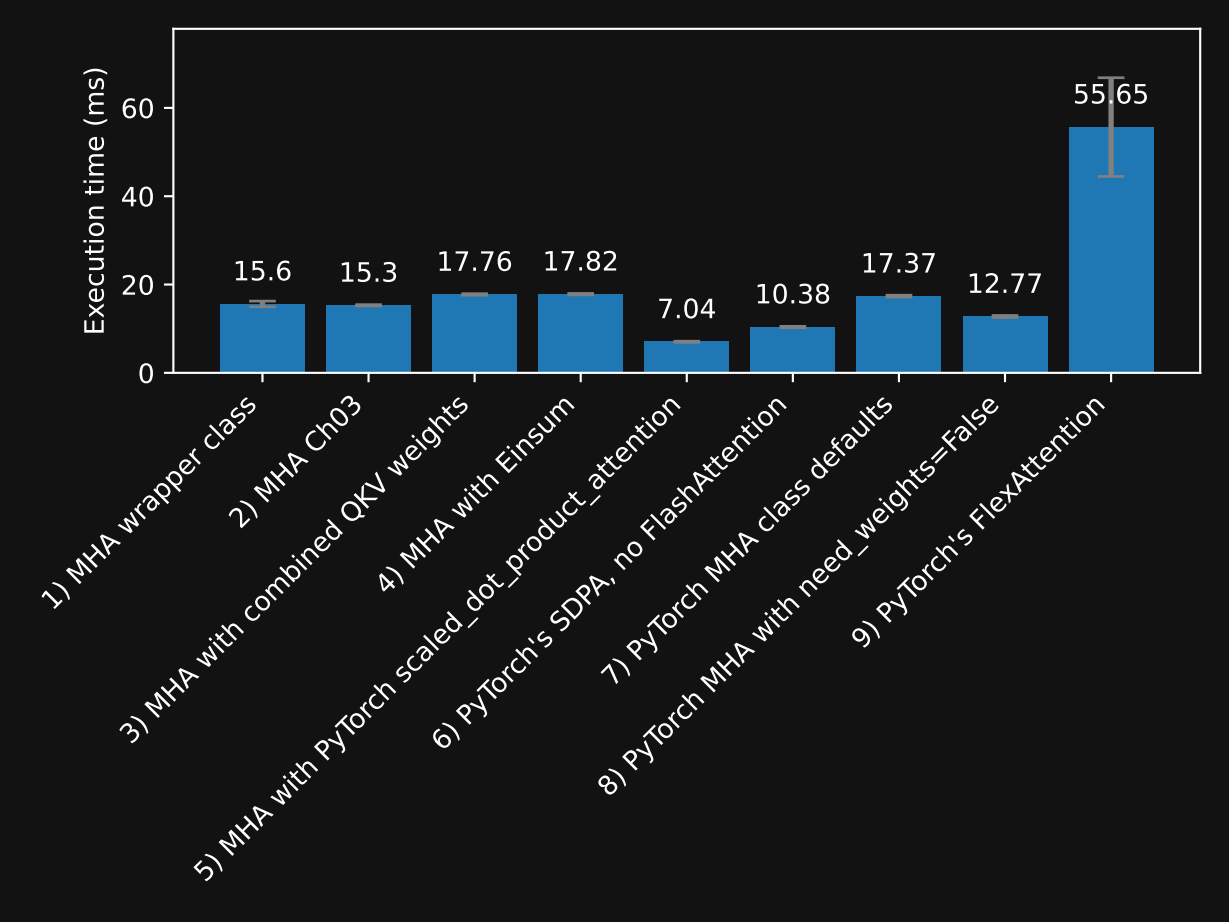
Execution time (ms)
0
20
40
60
15.6
15.3
17.76
17.82
7.04
10.38
17.37
12.77
55.65
1) MHA wrapper class
2) MHA Ch03
3) MHA with combined QKV weights
4) MHA with Einsum
5) MHA with PyTorch scaled_dot_product_attention
6) PyTorch's SDPA, no FlashAttention
7) PyTorch MHA class defaults
8) PyTorch MHA with need_weights=False
9) PyTorch's FlexAttention