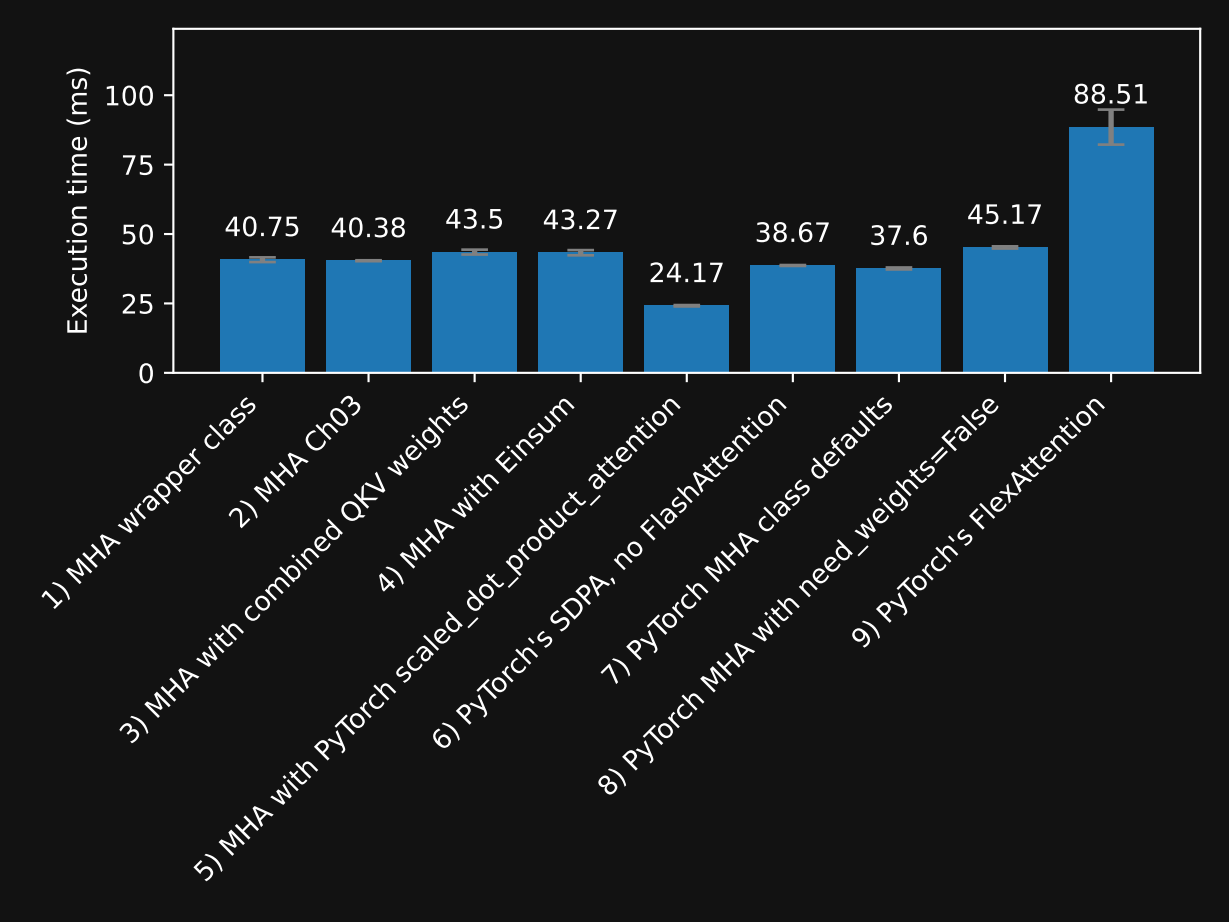
Execution time (ms)
0
25
50
75
100
40.75
40.38
43.5
43.27
24.17
38.67
37.6
45.17
88.51
1) MHA wrapper class
2) MHA Ch03
3) MHA with combined QKV weights
4) MHA with Einsum
5) MHA with PyTorch scaled_dot_product_attention
6) PyTorch's SDPA, no FlashAttention
7) PyTorch MHA class defaults
8) PyTorch MHA with need_weights=False
9) PyTorch's FlexAttention